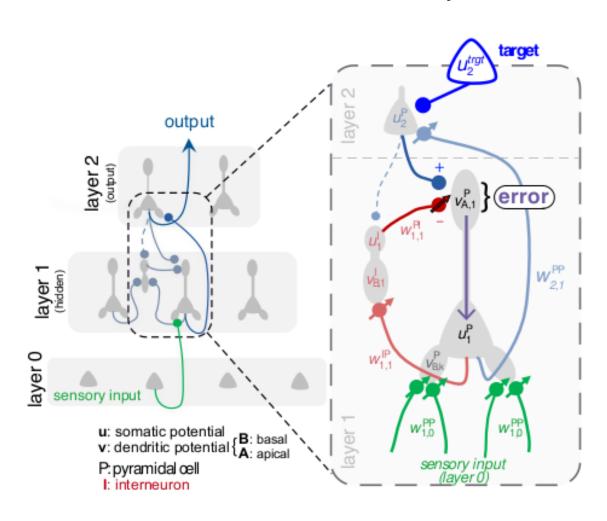
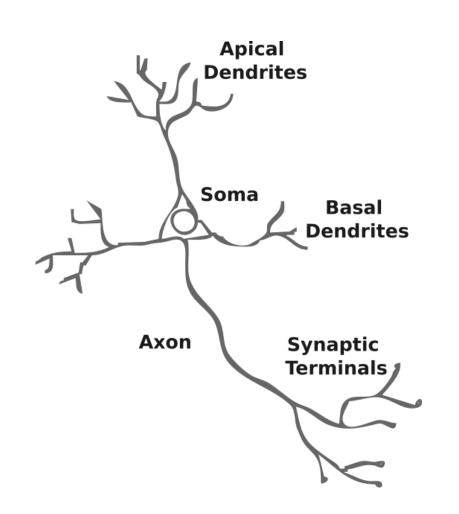
Dendritic cortical microcircuits approximate the backprob algorithm (Sacramento 2017)

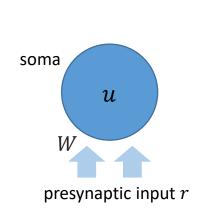


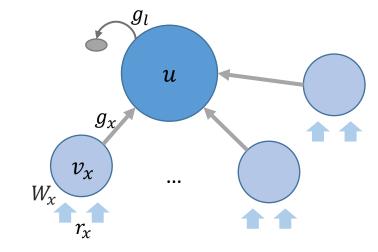
- reproduce lost code (→ repository with source and examples will be available online)
- How does the model work? Any subtleties?

(future: bring to Hardware)

From point neurons to multi-compartment neurons







$$\dot{u} = -g_l u + \sum_{x} g_x (v_x - u) \qquad \begin{aligned} v_x &= \sum_{i} W_{xi} r_{xi} \\ r &= \phi(u) \end{aligned}$$

$$\tau \dot{u} = -u + \tau \sum_{x} g_{x} v_{x}$$

$$1/\tau = g_{tot}$$

$$= g_{l} + \sum_{x} g_{x}$$

$$T\gg \tau$$

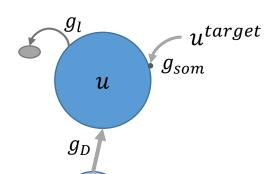
$$\tilde{u} = \tau \sum_{x} g_{x} v_{x}$$

steady-state solution

Learning by dendritic prediction of somatic spiking







$$\hat{v}_B \equiv \tilde{u} = \frac{g_D}{g_l + g_D} v_D = \frac{g_D}{g_l + g_D} W \cdot r$$

with nudging
$$\tilde{u} = (1 - \lambda)\hat{v}_B + \lambda u^{target}$$
, $\lambda = \frac{g_{som}}{g_{som} + g_{som}}$

$$\lambda = \frac{g_{som}}{g_l + g_D + g_{som}}$$

dendritic prediction error: $e \equiv \phi(u) - \phi(\hat{v}_B)$



$$\dot{W} = \eta \ e \ r^T$$



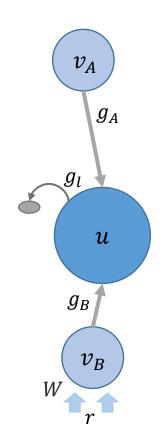
$$\dot{W} = \eta \ e \ r^T$$
 until $\hat{v}_B = u^{target}$

in praxis:
$$\tau_W \dot{\Delta} = -\Delta + e \ r^T \\ \dot{W} = \eta \Delta$$

Supplements; Urbanczik 2014

Pyramidal neurons





$$\hat{v}_B \equiv \tilde{u} = \frac{g_B}{g_l + g_B + g_A} v_B = \frac{g_B}{g_l + g_B + g_A} W \cdot r$$

$$v_A \approx 0$$

$$\tilde{u} = \hat{v}_B + \lambda v_A, \quad \lambda = \frac{g_A}{g_l + g_B + g_A}$$

dendritic prediction error: $e \equiv \phi(u) - \phi(\hat{v}_B)$



$$\dot{W} = \eta \ e \ r^T$$

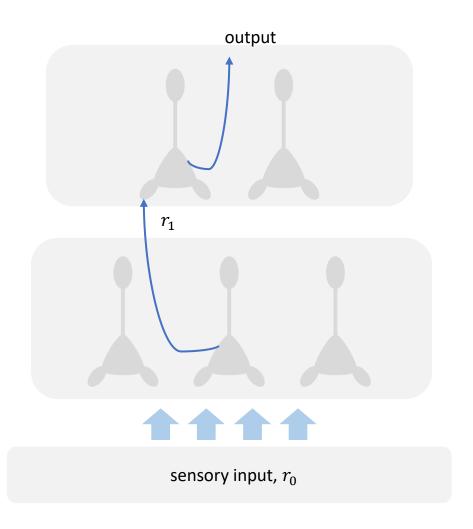


until forever?

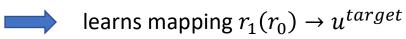


 $V_{\mathsf{B}}^{\mathsf{P}}$

Deep cortical microcircuits (naïve)



apply the learning rule in the output layer

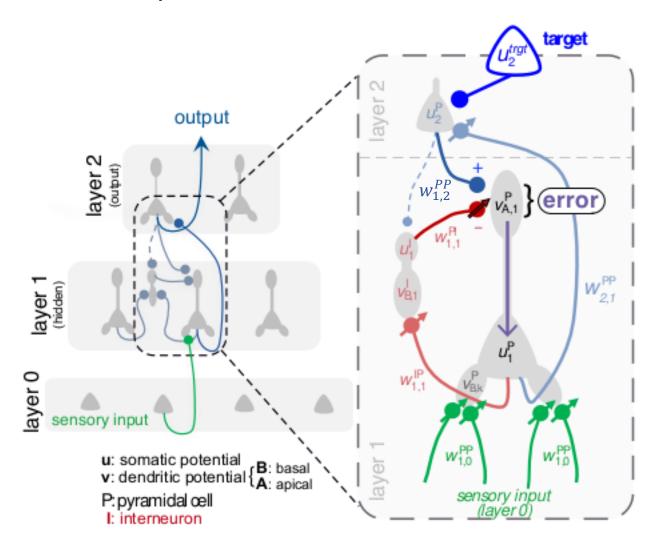


How can the hidden layer support learning? How should it adjust $r_1(r_0)$?

propagate output error back to hidden layer via $r_1(r_0)$



Deep cortical microcircuits



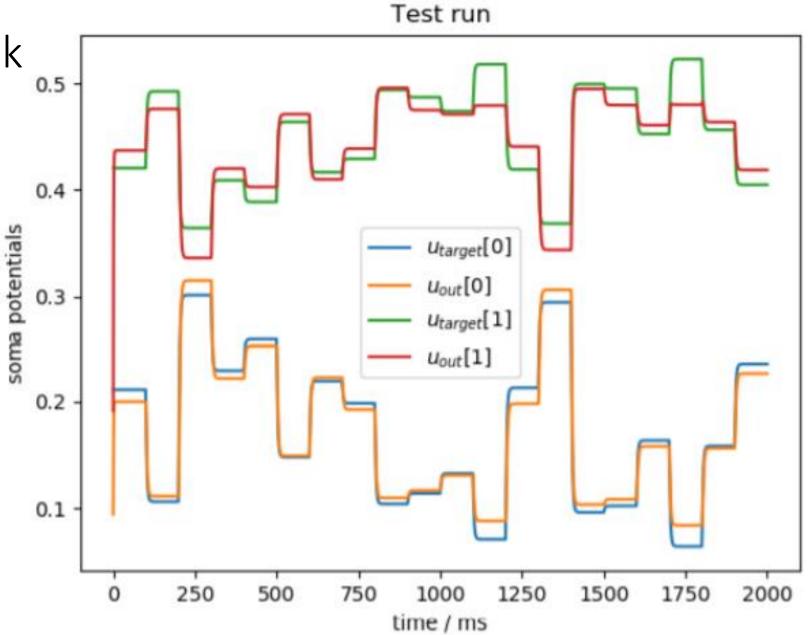
lateral plasticity → self-predicting state:

$$W_{2,1}^{PP} = W_{1,1}^{IP} \longrightarrow v_A^P \approx 0$$

$$W_{1,1}^{PI} = -W_{1,2}^{PP} \longrightarrow u_1^I \approx u_2^I$$

Mimic-regression-task

task: 2x3x2 network learns to mimic the output of a randomly initialized network of same size



Steady-State Approximation

100 ms pattern time \gg 0.1 ms time step dt \rightarrow training is slow idea: approximate the settling process

- 1. propagate the input upwards, layer by layer with dendritic prediction / steady-state solution:
- 2. propagate corrections downwards with full steady-state solution:
- 3. repeat n_{passes} times \rightarrow (hopefully) converge to steady-state of entire network

Plasticity:

- solve $\dot{W} = \eta \ e \ r^T$ with time step $dT \gg dt$
- update potentials after each weight update

What do we hope to gain? Approximation might not be accurate, but good approach for hyperparameter search!

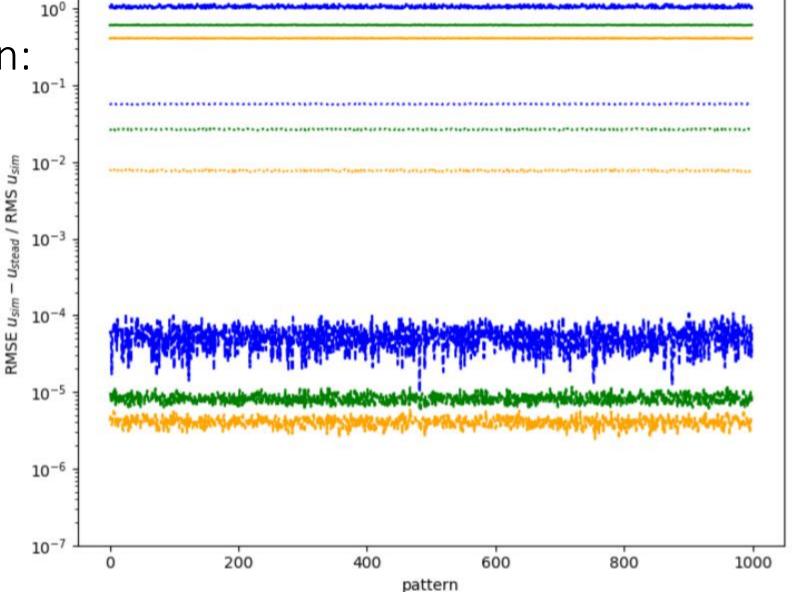


Steady-State Approximation: Validity

relative rms deviation of

soma potentials, when

plasticity is switched off

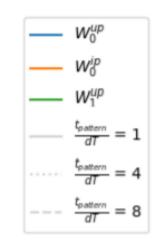


n_passes: 1 n_passes: 2 n_passes: 4

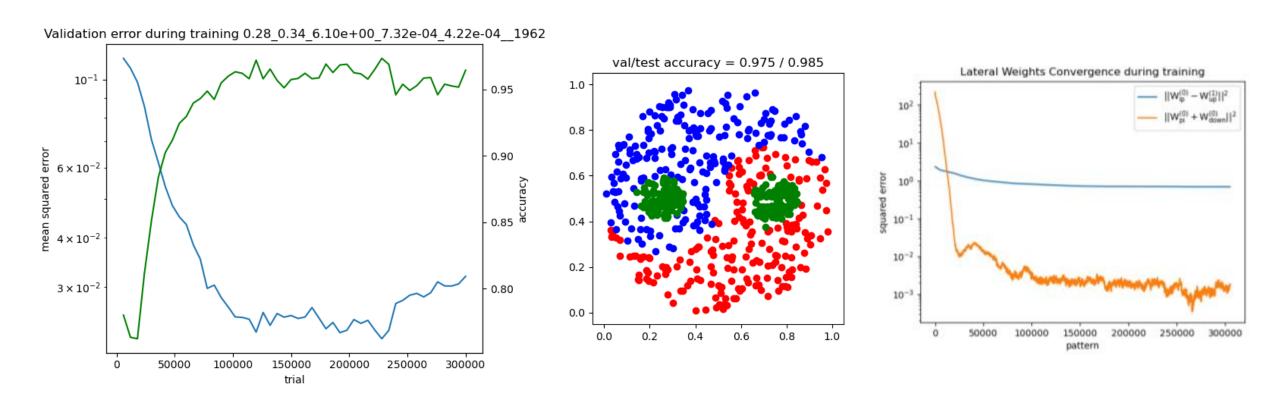
Simulation vs. Steady-State: Divergence of weights

pattern

Steady-State Approximation: Validity Wstead / RMS Wsim 10^{-1} RMSE Wsim 10^{-2} 10^{-3} relative rms deviation of weights 10^{-4} 1000 2000 3000 4000 5000



Yin-Yang results



trained on 55 epochs a 6000 samples: $(97.2 \pm 1.1)\%$ accuracy

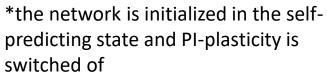
Yinyang PyraLNet performance vs learning-lag

Yin-Yang results: learning-lag

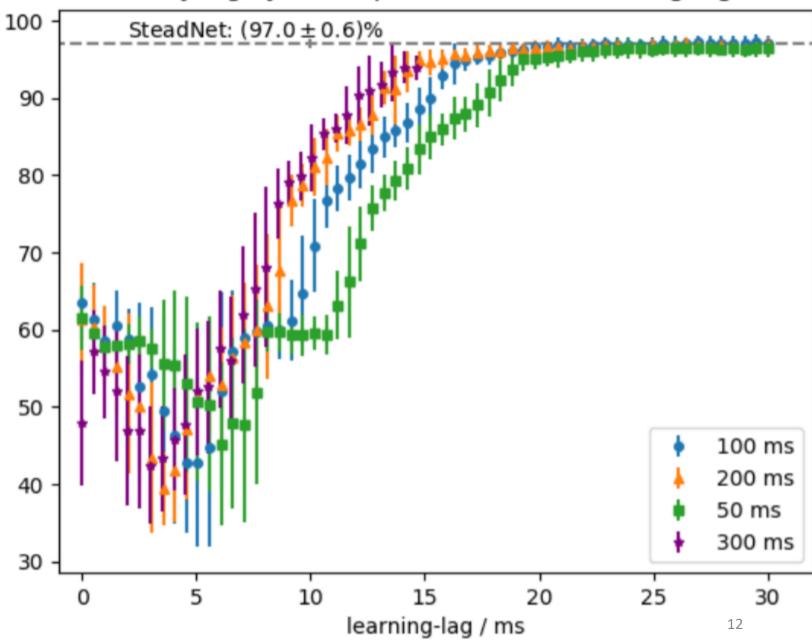
during settling unwanted plasticity is induced → chokes learning beyond a certain point

accuracy / %

introduce learning-lag: no plasticity for fixed time after each new pattern



^{*}trained on 6000 samples for 45 epochs



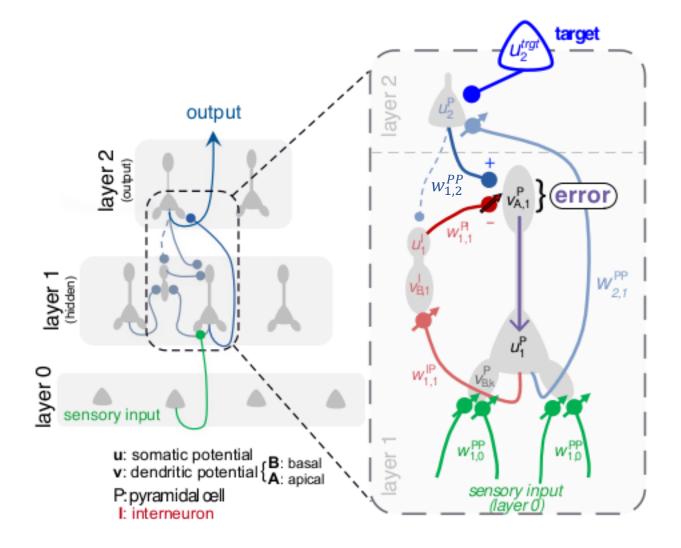
Summary

- Code for simulation (PyraLNet) and steady-state approximation (SteadNet) is available
- model is able to solve the Yin-Yang problem effectively once learning-lag is introduced
- SteadNet is able to solve MNIST

However:

 SteadNet results cannot directly be generalized to full simulation (since unwanted plasticity during settling might cause problems)

Deep cortical microcircuits

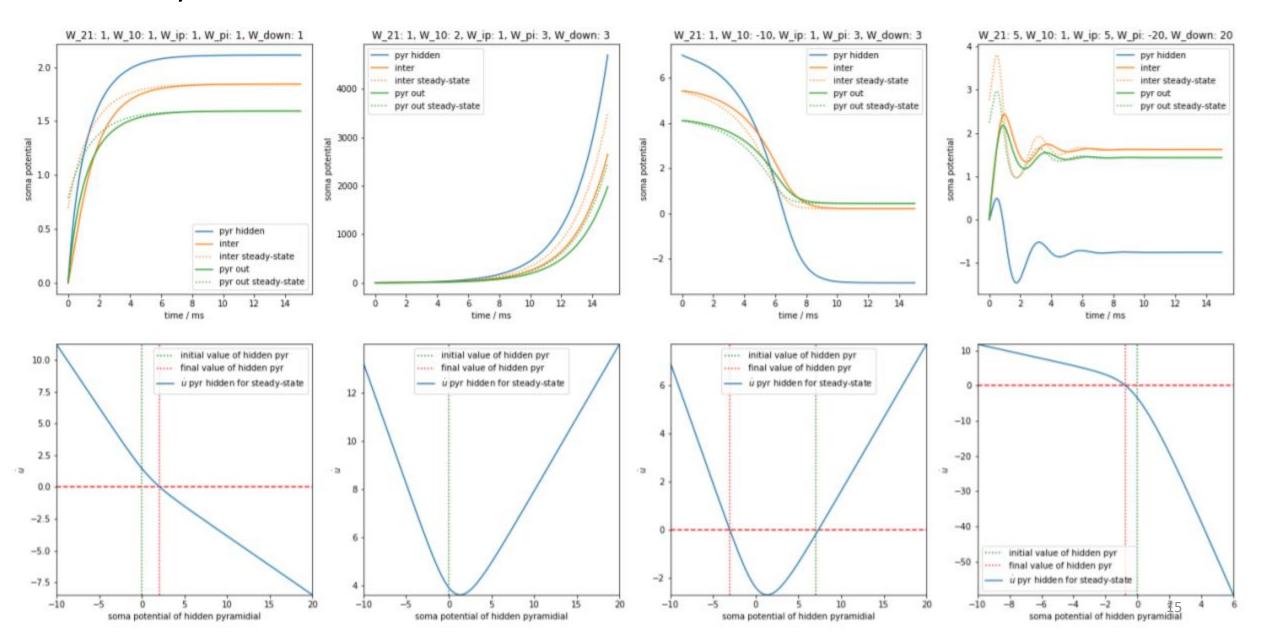


$$\begin{split} u_{2}^{P} &= -g_{tot}u_{2}^{P} + \mathsf{g_{B}}W_{2,1}^{PP}\phi(u_{1}^{P}) + g_{som}u^{target} \\ u_{1}^{I} &= -g_{tot}u_{1}^{I} + \mathsf{g_{D}}W_{1,1}^{IP}\phi(u_{1}^{P}) + g_{som}u_{2}^{P} \\ u_{1}^{P} &= -g_{tot}u_{1}^{P} + \mathsf{g_{B}}W_{1,0}^{PP}r_{0} + g_{A}v_{A,1}^{P} \\ v_{A,1}^{P} &= W_{1,1}^{PI}\phi(u_{1}^{I}) + W_{1,2}^{PP}\phi(u_{2}^{P}) \\ \dot{W}_{1,1}^{PI} &= -\eta^{PI}v_{A,1}^{P} \; , \quad \dot{W}_{1,2}^{PP} &= 0 \end{split}$$

lateral plasticity → self-predicting state:

$$W_{2,1}^{PP} = W_{1,1}^{IP}$$
 $W_{1,1}^{PI} = W_{1,2}^{PP}$
 $v_A^P \approx 0$
 $u_1^I \approx u_2^P$

Stability

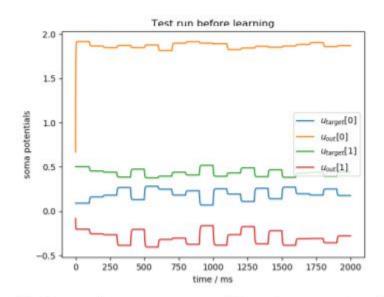


Mimic-regression-task

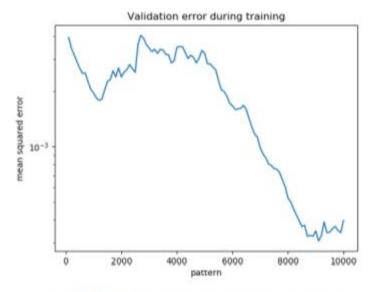
task: 2x3x2 network learns to mimic the output of a randomly initialized network of same size

Details:

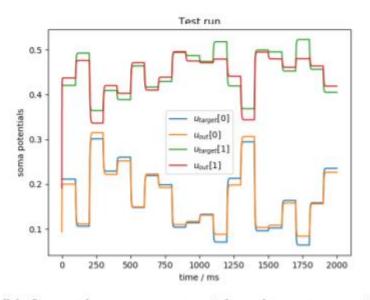
- g = 0.1 1, $\tau \approx 0.5$ ms
- weights of order 1 (elements)
- input, output patterns are presented for 100 ms and are smoothly transitioned in between (3 ms)
- plasticity time constant $\tau_W = 30 \text{ ms}$
- simulation time step: 0.1 ms



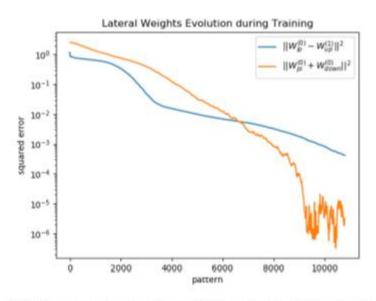
(a) Output layer soma potentials and target potentials before training.



(c) MSE validation error during training.

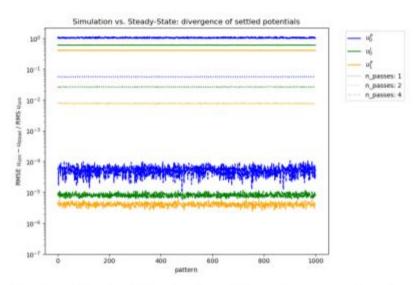


(b) Output layer soma potentials and target potentials after training.

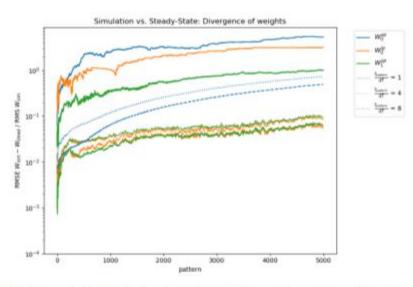


(d) Emergence of self-predicting-state during training starting from randomly initialized confections.

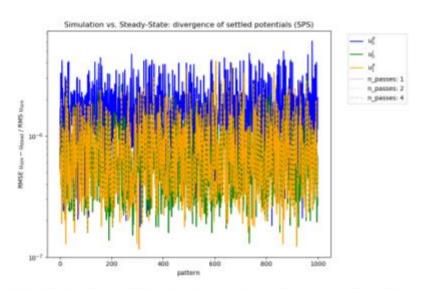
Steady-State Approximation: Validity



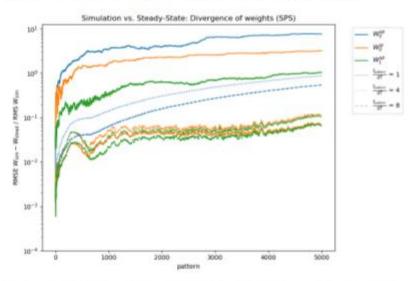
(a) Deviation⁵ of the simulated from the approximation soma potentials, when plasticity is switched off.



(c) Deviation of the simulated from the approximation weights, when plasticity is switched on.



(b) Deviation of the simulated from the approximation soma potentials, when plasticity is switched off, but the network is initialized in the self-predicting state.

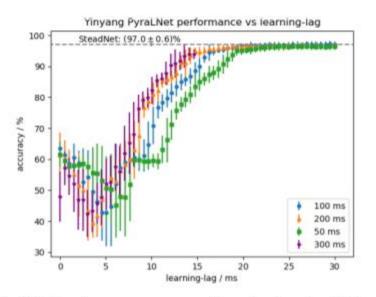


(d) Deviation of the simulated from the approximation weights, when plasticity is switched on, but the network is initialized in the self-predicting state.

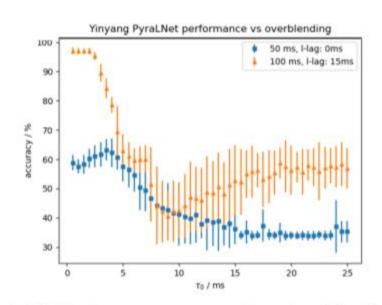
Yin-Yang results: learning-lag

during settling unwanted plasticity is induced → chokes learning beyond a certain point

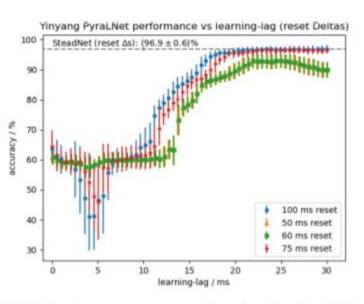
introduce learning-lag: no plasticity for fixed time after each new pattern



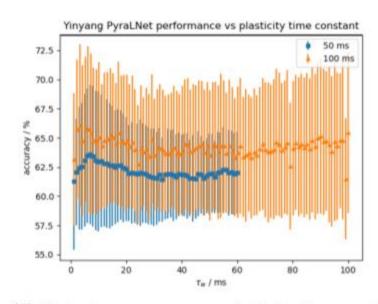
(a) Test set accuracy versus learning-lag for different $t_{pattern}$. See main text for parameter details.



(c) Test set accuracy versus pattern transition time constant τ_0 .



(b) Test set accuracy versus learning-lag for different t_{pattern}, but with reset_deltas switched on.



(d) Test set accuracy versus plasticity time constant τ_w.

^{*}the network is initialized in the selfpredicting state and PI-plasticity is switched of

^{*}trained on 6000 samples for 45 epochs